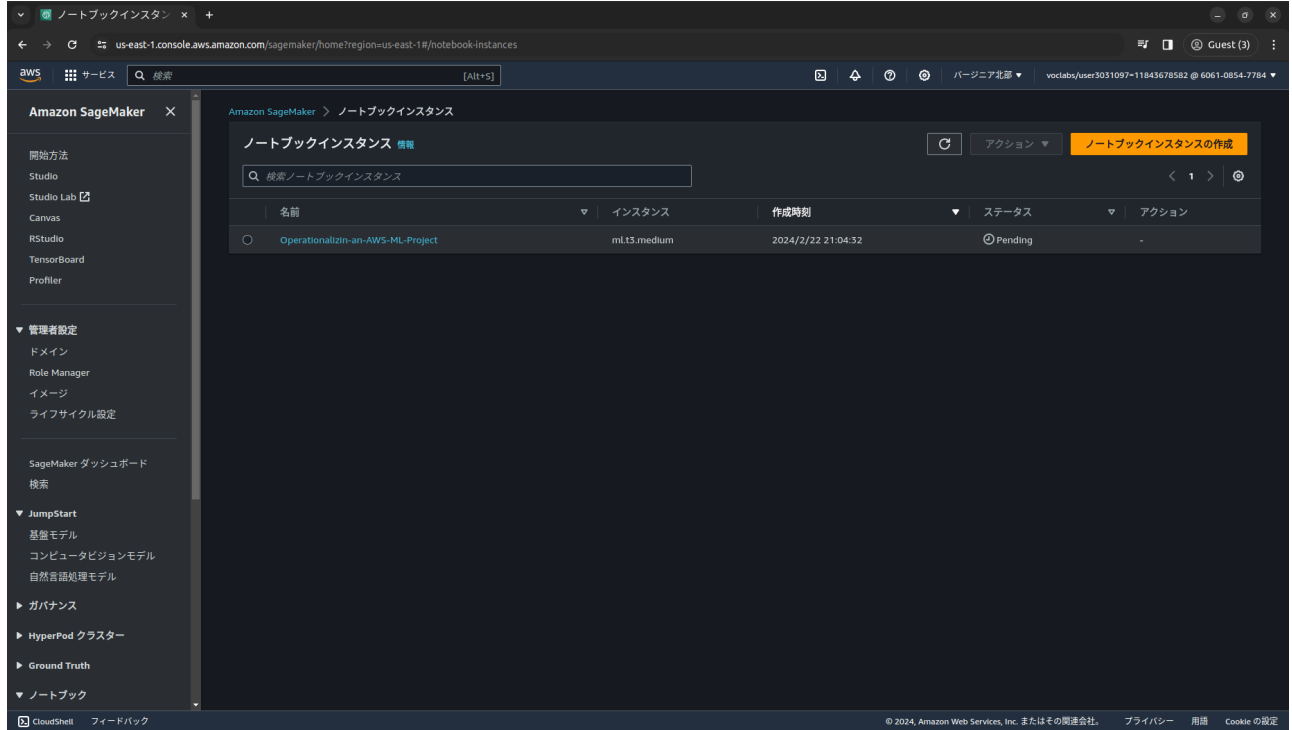


Project: Operationalizing an AWS ML Project

Training and deployment

- Student sets up a Sagemaker instance:



- Student justifies the type of Sagemaker instance they created: **ml.t3.medium** was created as it is sufficient to run my notebook.

Reasons for Selecting Sagemaker Instance "ml.t3.medium"

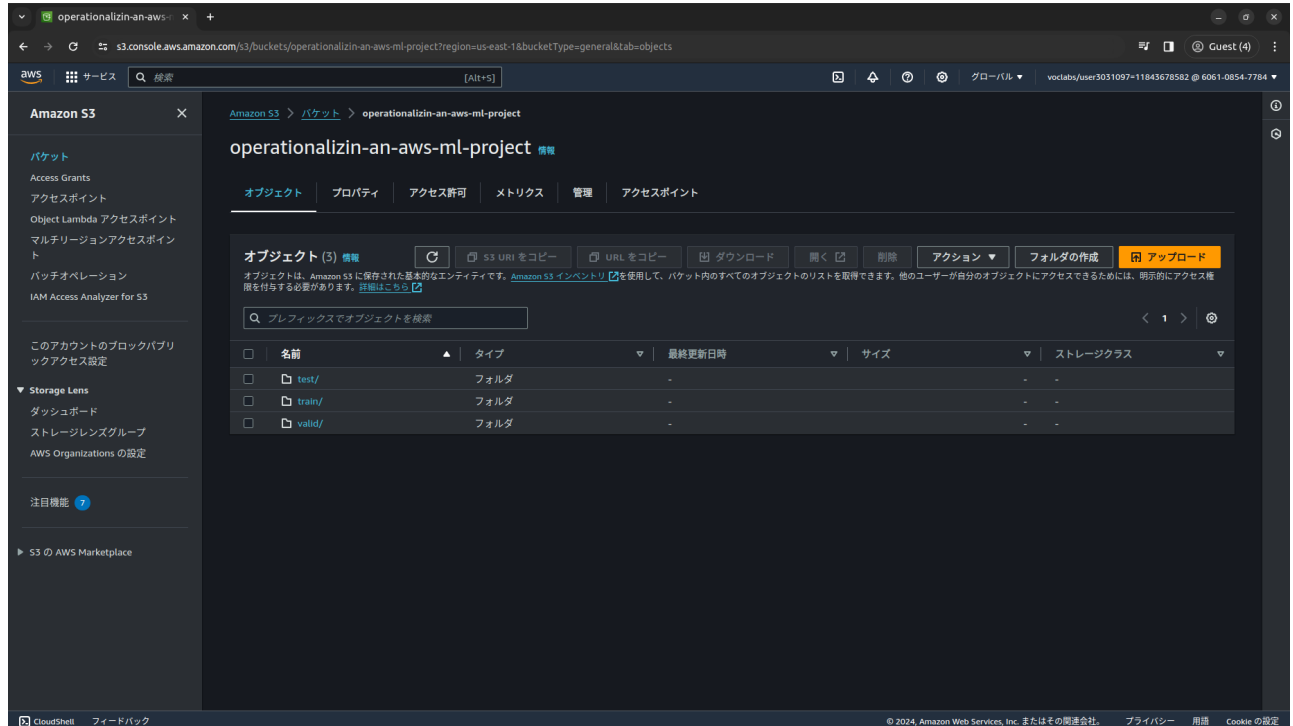
When choosing the Sagemaker instance type "ml.t3.medium," several factors need to be considered, including power, speed, and cost efficiency.

- **Power** : The "ml.t3.medium" instance offers sufficient computational power for many common machine learning tasks. While it may not be the most powerful option available, it provides adequate resources for training and deploying models on small to medium-sized datasets.
- **Speed** : In terms of speed, the "ml.t3.medium" instance strikes a balance between performance and cost. It may not be the fastest instance type available, but it still offers reasonable processing speeds for training and inference tasks. For smaller datasets and less computationally intensive workloads, it can deliver satisfactory performance without incurring excessive costs.
- **Cost Efficiency** : One of the primary reasons for choosing the "ml.t3.medium" instance is its cost efficiency. As a mid-tier instance type, it typically offers a lower hourly rate compared to higher-powered instances like "ml.m5.large" or "ml.c5.xlarge." This makes it an attractive option for users who prioritize cost savings without compromising too much on performance.


- **Flexibility** : The "ml.t3.medium" instance provides a good balance of CPU and memory resources, making it suitable for a wide range of machine learning tasks. Whether you're training supervised learning models, performing data preprocessing, or deploying inference endpoints, this instance type offers the flexibility to handle various workloads effectively.

In summary, selecting the "ml.t3.medium" instance type on Sagemaker can be justified by its adequate power, reasonable speed, cost efficiency, and flexibility, making it a suitable choice for many machine learning projects, especially those with budget constraints or moderate computational requirements.

- Student creates an S3 bucket for data downloads:



- Student deploys an endpoint by running all starter notebooks:



- Student accomplishes multi-instance training: [Jupyter-Notebook](#)

EC2 Training

- Create and open an EC2 instance: **m5.xlarge** with **Deep Learning AMI GPU PyTorch 2.1.0** was selected due to its low cost.

```

aws
[us-east-1.console.aws.amazon.com/ec2-instance-connect/ssh?region=us-east-1&connType=standard&instanceId=i-06fbcd352040ef402&osUser=root&sshPort=22#]

[root@ip-172-31-18-61 ~]# python solution.py
/opt/conda/lib/python3.10/site-packages/torchvision/models/_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the future, please use 'weights' instead.
  warnings.warn(
/opt/conda/lib/python3.10/site-packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a weight enum or 'None' for 'weights' are deprecated since 0.13 and may be removed in the future. The current behavior is equivalent to passing 'weights=ResNet50_Weights.IMAGENET1K_V1'. You can also use 'weights=ResNet50_Weights.DEFAULT' to get the most up-to-date weights.
  warnings.warn(msg)
Downloading: "https://download.pytorch.org/models/resnet50-0676ba61.pth" to /root/.cache/torch/hub/checkpoints/resnet50-0676ba61.pth
100% | 97.8M/97.8M [00:00<00:00, 202MB/s]
Starting Model Training
saved
[root@ip-172-31-18-61 ~]# ll
total 1108560
drwxr-xr-x 5 root root    44 Mar 27  2017 dogImages
-rw-r--r-- 1 root root 1132023110 Apr  1  2017 dogImages.zip
-rw-r--r-- 1 root root    4843 Feb 22 13:53 solution.py
drwxr-xr-x 2 root root    23 Feb 22 14:05 TrainedModels
[root@ip-172-31-18-61 ~]# ll TrainedModels/
total 93212
-rw-r--r-- 1 root root 95445794 Feb 22 14:05 model.pth
[root@ip-172-31-18-61 ~]#

```

Reasons for Selecting EC2 Instance "m5.xlarge"

When opting for the EC2 instance type "m5.xlarge," several factors come into play, including power, speed, and cost efficiency.

- **Power** : The "m5.xlarge" instance offers significant computational power suitable for a wide range of compute-intensive tasks. With its combination of vCPUs, memory, and network performance, it can handle demanding workloads such as data processing, machine learning training, and high-performance computing tasks effectively.
- **Speed** : In terms of speed, the "m5.xlarge" instance provides excellent performance for compute-bound applications. It features high-frequency Intel Xeon processors and enhanced networking capabilities, enabling faster data transfer and processing. This translates to reduced processing times and improved overall efficiency for time-sensitive tasks.
- **Cost Efficiency** : While the "m5.xlarge" instance may not be the most economical option available, it offers a good balance between performance and cost. Its hourly rate reflects the higher computational power and enhanced features compared to lower-tier instances. However, for workloads that demand superior performance, the increased cost may be justified by the time saved in completing tasks more quickly.
- **Resource Availability** : The "m5.xlarge" instance provides ample resources, including CPU cores, memory, and storage, making it suitable for a wide range of applications. Whether you're running data analytics, hosting web applications, or deploying machine learning models, this instance type offers the necessary resources to support your workload efficiently.

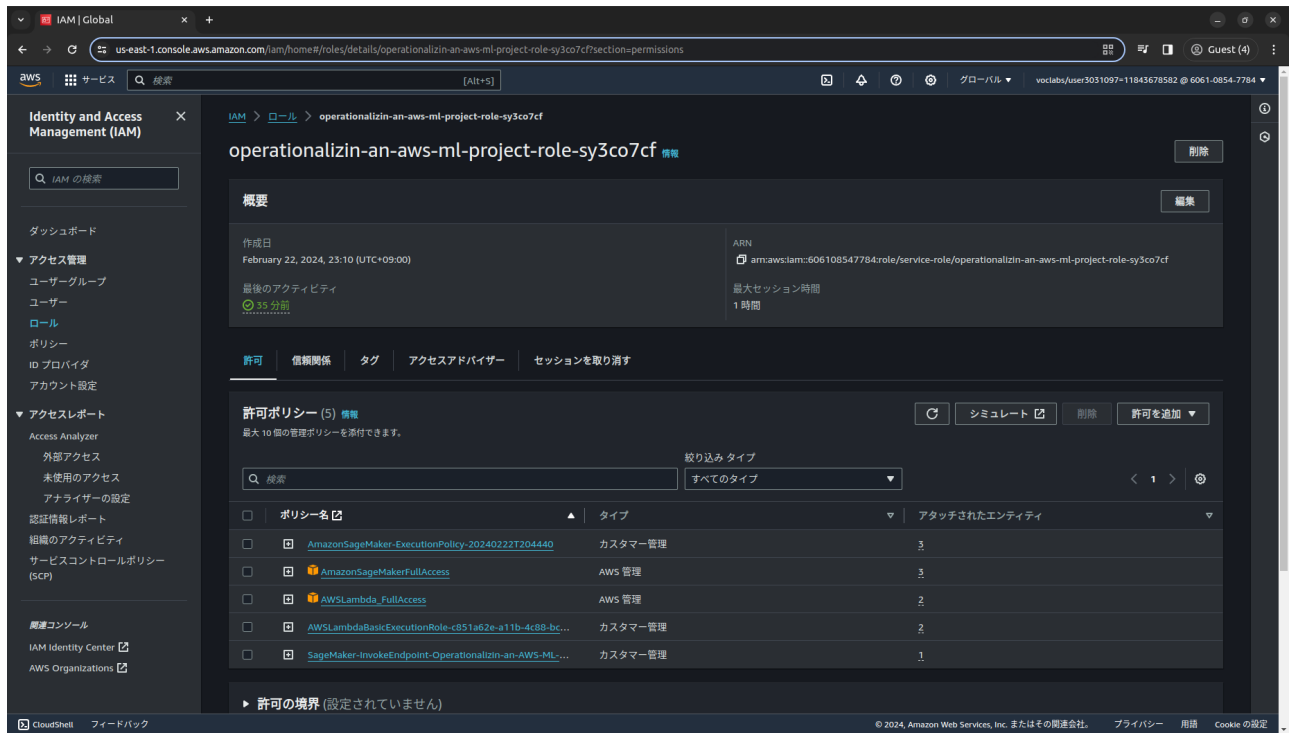
- In summary, selecting the "m5.xlarge" EC2 instance offers significant power, high-speed performance, reasonable cost efficiency, ample resource availability, and reliable operation, making it an excellent choice for compute-intensive applications where performance is paramount.

- ## Lambda function setup

-

- ## Security and testing

- Attach a security policy to the role associated with the Lambda endpoint :



- Test the Lambda function by invoking the endpoint with the "lab.jpg" input :

```
{
  "statusCode": 200,
  "headers": {
    "Content-Type": "text/plain",
    "Access-Control-Allow-Origin": "*"
  },
  "type-result": "<class 'str'>",
  "Content-Type-In": "LambdaContext([aws_request_id=8ce7215c-f463-4003-a8e8-2eb52971c02e, log_group_name=/aws/lambda/project-4, log_stream_name=2023/04/12/[$LATEST]779f63f6e8f74c13a51970470403d45d, function_name=project-4, memory_limit_in_mb=128, function_version=$LATEST, invoked_function_arn=arn:aws:lambda:us-east-1:271564095025:function:project-4, client_context=None, identity=CognitoIdentity([cognito_identity_id=None, cognito_identity_pool_id=None]))",
  "body": "[[-1.4969292879104614, -2.295322895050049, 1.7711918354034424, 1.7000603675842285, -2.9478816986083984, -2.8578217029571533, -4.275012493133545, -0.1858823001384735, -7.876079559326172, 3.485708713531494, 1.0156581401824951, -5.376102924346924, 0.08712732791900635, 1.043073058128357, -6.567814350128174, -4.29003381729126, -6.980752468109131, 0.8602855205535889, -1.9477488994598389, 3.517031192779541, 1.1386747360229492, 2.9857466220855713, -7.429625034332275, -4.872533321380615, -6.990846633911133, -6.551341533660889, -2.713444471359253, -7.38569974899292, -4.197381973266602, 0.5921437740325928, -0.23152270913124084, -1.7627674341201782, -4.5388617515563965, 0.545304000377655, -3.987135171890259, -3.03073787689209, -3.0044939517974854, -0.24150973558425903, 1.3483713865280151, -1.685007095336914, -1.484816312789917,
```

```

1.8437116146087646, 3.838435649871826, 0.6947075128555298,
-0.1137256920337677, -11.679482460021973, 1.3911057710647583,
-1.805575966835022, -2.34354567527771, 1.7085063457489014,
-0.9614053964614868, -8.280823707580566, -7.071294784545898,
0.16166920959949493, -0.5199874639511108, -0.843184232711792,
-6.246165752410889, -2.350785732269287, -1.3200641870498657,
-2.318167209625244, -4.5581793785095215, -8.424875259399414,
-7.389919281005859, -8.225934028625488, -5.960506916046143,
-6.004213333129883, 3.052112102508545, -0.9716691374778748,
0.1721428632736206, 1.1353175640106201, 5.675631046295166,
-4.881450176239014, -4.949894428253174, -4.2738471031188965,
-1.0440956354141235, 0.615058183670044, -7.484931468963623,
-3.1543450355529785, -4.828455448150635, -7.795574188232422,
2.6574225425720215, -8.622648239135742, 1.655287504196167,
1.2484384775161743, -8.302145957946777, -3.6932826042175293,
2.7479288578033447, -7.128697395324707, 1.1833446025848389,
2.0749804973602295, -8.263742446899414, -2.6031601428985596,
-2.9580190181732178, -6.628727436065674, -2.2038745880126953,
0.8574008941650391, 0.36140191555023193, 1.240427851676941,
-6.453752517700195, -9.306218147277832, -5.975752353668213,
-0.851379930973053, -3.8800625801086426, -4.767630577087402,
-2.8386266231536865, -5.776670932769775, -1.417716383934021,
2.4901983737945557, 1.7963000535964966, 0.9240027666091919,
1.3527525663375854, 1.043146014213562, -5.193169116973877,
-2.356482982635498, -7.384207248687744, -0.1713782399892807,
-4.5393853187561035, -2.649200916290283, -4.741747856140137,
0.618251621723175, 0.2919156551361084, -3.7748732566833496,
-3.141923666000366, -1.8757330179214478, -6.020362377166748,
-5.679347991943359, -1.4780714511871338, 1.1916701793670654,
-4.358892917633057, -7.019332408905029, -5.486293792724609,
3.572416067123413, -3.9262099266052246]]]"
}

```

- Student identified potential vulnerabilities in the project's IAM configuration :

Roles with "FullAccess" Policies: Roles that are granted "FullAccess" policies pose a significant security risk as they grant unrestricted access to resources and services within the system. These roles essentially provide overly broad permissions, which can lead to unauthorized access, data breaches, or even system compromise if exploited by malicious actors.

Concurrency and auto-scaling

Student clearly describes the configuration of concurrency (on the Lambda function) and auto-scaling (for the deployed endpoint)

In AWS, effective concurrency and auto-scaling configurations are vital for performance and cost efficiency. This overview delves into configuring both for Lambda functions and deployed endpoints, alongside traffic, cost, and efficiency considerations.

- **Concurrency (Lambda Function) :**
 - Set maximum simultaneous executions.

- Use AWS Console or CLI.
- Consider reserved concurrency for precise control.
- **Auto-scaling (Deployed Endpoint) :**
 - Utilize services like EC2 Auto Scaling or Elastic Beanstalk.
 - Define scaling policies and integrate with CloudWatch.
 - Considerations:
- **Traffic :**
 - Analyze patterns and adjust configurations accordingly.
- **Cost :**
 - Balance performance with cost by selecting appropriate instances and optimizing configurations.
- **Efficiency :**
 - Optimize resource usage through controls like serverless computing and load balancing.

Optimizing concurrency and auto-scaling in AWS enhances performance, saves costs, and maximizes resource utilization. Regular monitoring and adjustment ensure continued efficiency.